

Lecture 22: The Singular Value Decomposition (SVD)

This lecture introduces the *singular value decomposition*, which we will use to compress images.

SVD

The SVD is a decomposition of an $n \times m$ matrix A as

$$A = U\Sigma V^T$$

where $U \in \mathbb{R}^{n \times n}$ and $V \in \mathbb{R}^{m \times m}$ are orthogonal and $\Sigma \in \mathbb{R}^{n \times m}$ is diagonal, of the form ($n \geq m$)

$$\Sigma = \begin{pmatrix} \sigma_1 & & & \\ & \sigma_2 & & \\ & & \ddots & \\ & & & \sigma_m \\ & & & 0 \\ & & & \vdots \\ & & & 0 \end{pmatrix},$$

where $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_m \geq 0$.

The SVD can be calculated using the `svd` command:

In [68]:

```
A=rand(5,5)
U,σ,V=svd(A)

Σ=diagm(σ)

norm(A-U*Σ*V')
```

Out[68]:

```
2.0261379083813132e-15
```

U and V are indeed orthogonal:

In [70]:

```
norm(U'*U-I), norm(V'*V-I)
```

Out[70]:

```
(1.3615860578367606e-15, 6.210964166796445e-16)
```

Relationship with eigenvalue decomposition

The SVD is closely connected to the eigenvalue decomposition for symmetric matrices. The difference is that eigenvalues are signed while singular values are all positive:

In [2]:

```
A=rand(5,5)
A=A+A'          # any matrix plus its transpose is symmetric
svdvals(A), eigvals(A) # the magnitudes are the same, but the sign and order
ing are different
```

Out[2]:

```
([5.169732137294604, 1.6312252622180625, 1.0174977889866161, 0.89768
29470592561, 0.7363935801508221], [-1.0174977889866166, -0.897682947
0592556, 0.7363935801508223, 1.6312252622180614, 5.169732137294606])
```

A symmetric matrix has an eigenvalue decomposition

$$A = Q\Lambda Q^T$$

where Q is orthogonal (we'll explain this in two lectures):

In [3]:

```
λ, Q=eig(A)
Λ=diagm(λ)
norm(A-Q*Λ*Q')
```

Out[3]:

```
5.706584266746393e-15
```

For symmetric matrices, singular values are the absolute value of eigenvalues. We can see this by constructing the SVD from an eigenvalue decomposition by permuting and taking the absolute value:

In [5]:

```
P=eye(5)[ :,[5,4,1,3,2]]  
  
norm(A-Q*P*P'*abs(Λ)*P*P'*sign(Λ)*Q')  
  
Ū=Q*P  
Σ̃=P'*abs(Λ)*P  
V̄=Q*sign(Λ)*P  
  
norm(Ū*Σ̃*V̄'-A)
```

Out[5]:

5.632251628873568e-15

Singular values are not (necessarily) absolute value of eigenvalues for non-symmetric matrices: in the following matrix, the eigenvalues are zero but the singular values are not:

In [6]:

```
A=[0 1.0;  
   0 0]  
  
svdvals(A) # returns σ, the list of singular values
```

Out[6]:

```
2-element Array{Float64,1}:  
 1.0  
 0.0
```

Properties of the SVD:

1) The 2-norm of a matrix is the largest singular value:

$$\|A\|_2 = \sigma_1 :$$

In [77]:

```
A=rand(5,3)  
U,σ,V=svd(A)  
maximum(σ),norm(A)
```

Out[77]:

(2.0999467726097785,2.099946772609779)

This follows since orthogonal matrices do not change 2-norms:

$$\|A\|_2 \triangleq \sup_{\|\mathbf{w}\|=1} \|A\mathbf{w}\| = \sup_{\|\mathbf{w}\|=1} \|U\Sigma V^\top \mathbf{w}\| = \sup_{\|\mathbf{w}\|=1} \|U\Sigma V^\top \mathbf{w}\| = \sup_{\|\mathbf{v}\|=1} \|\Sigma\| = \max_k \sigma_k = \sigma_1$$

2) The rank of a matrix is the number of non-zero singular values

This follows since the rank of a diagonal matrix (in this case Σ) is the number of non-zero entries, and U and V are of full rank.

3) The kernel of a rank- r matrix is the span of $\mathbf{v}[:, r+1:\text{end}]$.

This follows since:

$$A\mathbf{v}_k = U\Sigma V^\top V\mathbf{e}_k = U\Sigma\mathbf{e}_k = 0$$

if $k > r$.

The best rank- r approximation

Given the SVD, the best rank- r approximation is given by dropping all but the first r singular values:

$$A_r \triangleq U\Sigma_r V^\top$$

where

$$\Sigma_r \triangleq \begin{pmatrix} \sigma_1 & & & & & \\ & \sigma_2 & & & & \\ & & \ddots & & & \\ & & & \sigma_r & & \\ & & & & 0 & \\ & & & & & \ddots \\ & & & & & & 0 \\ & & & & & & \vdots \\ & & & & & & & 0 \end{pmatrix},$$

Theorem A_r is the best 2-norm approximation of A by a rank r matrix, that is, for all rank- r matrices B , we have

$$\|A - A_r\|_2 \leq \|A - B\|_2.$$

Proof We have

$$\|A - A_r\|_2 = \|U \begin{pmatrix} 0 & & & & & \\ & \ddots & & & & \\ & & 0 & & & \\ & & & \sigma_{r+1} & & \\ & & & & \ddots & \\ & & & & & \sigma_n \\ & & & & & \vdots \\ & & & & & 0 \end{pmatrix}\|_2 = \sigma_{r+1}$$

Suppose a rank- r matrix B has

$$\|A - B\|_2 < \|A - A_r\|_2 = \sigma_{r+1}.$$

For all $\mathbf{w} \in \ker(B)$ we have

$$\|A\mathbf{w}\|_2 = \|(A - B)\mathbf{w}\|_2 \leq \|A - B\| \|\mathbf{w}\|_2 < \sigma_{r+1} \|\mathbf{w}\|_2$$

But for all $\mathbf{w} \in \text{span } V[1 : r + 1]$, that is, $\mathbf{w} = V[1 : r + 1]\mathbf{c}$ for some $\mathbf{c} \in \mathbb{R}^{r+1}$ we have

$$\|A\mathbf{w}\|_2^2 = \|U\Sigma[:, 1 : r]\mathbf{c}\|_2^2 = \|\Sigma[:, 1 : r]\mathbf{c}\|_2^2 = \sum_{k=1}^{r+1} (\sigma_k c_k)^2 \geq \sigma_{r+1}^2 \|\mathbf{c}\|^2,$$

$$\text{i.e., } \|A\mathbf{w}\|_2^2 \geq \sigma_{r+1}^2 \|\mathbf{c}\|^2.$$

The dimension of the span of $\ker(B)$ is at least $n - r$, but the dimension of $\text{span } V[1 : r + 1]$ is at least $r + 1$. Since these two spaces cannot intersect we have a contradiction, since $(n - r) + (r + 1) = n + 1 > n$. ■

Application: image compression

We'll see an application of this to image compression: we are going to approximate an image A by its best rank- r approximation A_r . First load up a (greyscale) image:

In [78]:


```
using PyPlot

img=imread("/Users/solver/Desktop/ocean.png")
imshow(img)

R=img[:, :, 1]

A=1-R # swap white and black

imshow(A; cmap=:Greys);
```




The following gives the best rank-1 approximation:

In [79]:

```
U,σ,V=svd(A)

imshow(U[:,1]*σ[1]*V[:,1]'; cmap=:Greys);
```




Taking more and more singular values gives higher rank approximation:

In [80]:

```
U,σ,V=svd(A)


r=20
imshow(U[:,1:r]*diagm(σ[1:r])*V[:,1:r]';cmap=:Greys);
```



In [81]:

```
U,σ,V=svd(A)

r=50
imshow(U[:,1:r]*diagm(σ[1:r])*V[:,1:r]';cmap=:Greys);
```



We can store the resulting image using just the matrices $U[:,1:r]*\text{diagm}(\sigma[1:r])$ and $V[:,1:r]$, for a total of $rn + rm$ floating point numbers. For $r=20$, this is only about 20% of the storage as storing every pixel:

In [82]:

```
n,m=size(A)
(r*n+r*m)/(n*m)    # We use only 20% of the data as the full matrix
```

Out[82]:

0.20174507539100303